Neural Networks and Computational Intelligence
Practical Assignment II: Learning a linearly separable rule

You should either hand in a report for assignment II **or** assignment III. If you hand in assignment II **and** III, report III will be completely disregarded and not graded.

Learning a linearly separable rule
The topic of this assignment is the learning of a linearly separable rule from example data. We define outputs $S^\mu = \pm 1$ which are provided by a *teacher perceptron*. By construction, the resulting data set is guaranteed to be linearly separable, and learning in version space is a reasonable strategy in the absence of noise in the data set.

Consider a set of random input vectors as in assigment (I) with similar dimensions N. However, here we consider training labels S^μ which are defined as

$$S^\mu = \text{sign}(\mathbf{w}^* \cdot \boldsymbol{\xi}^\mu)$$

by a *teacher perceptron*. You should consider a randomly drawn $\mathbf{w}^*$ with $|\mathbf{w}^*|^2 = N$. Also, modify your code from assignment (I) so that it ...

- ... implements the sequential Minover algorithm:
  at each time step t, determine the stabilities

  $$\kappa^\nu(t) = \frac{\mathbf{w}(t) \cdot \boldsymbol{\xi}^\nu S^\nu}{|\mathbf{w}(t)|} \quad \text{for all examples } \nu$$

  and identify the example $\mu(t)$ that has currently the minimal stability $\kappa^{\mu(t)} = \min_\nu \{\kappa^\nu(t)\}$. In case of a *tie*, it does not matter which of the examples is chosen.
  With the selected example $\mu(t)$, perform a Hebbian update step

  $$\mathbf{w}(t+1) = \mathbf{w}(t) + \frac{1}{N} \boldsymbol{\xi}^{\mu(t)} S^{\mu(t)}$$

  and go to the next time step. In contrast to the Rosenblatt algorithm, the sequence of examples is not fixed and in each step a non-zero update is performed. Note that Minover should not stop when $\{E^\nu > 0\}_{\nu=1}^P$, because the stability will increase further. Run the algorithm until a number of $t_{max} = n_{max} \cdot P$ single training steps have been performed in total. The final weight vector $\mathbf{w}(t_{max})$ with stability $\kappa(t_{max})$ for a given set of data should approximate the perceptron of optimal stability $\mathbf{w}_{max}$. Initialize the weight vector as $\text{w}(t) = 0$ in each training process and implement a (reasonable) stopping criterion of your choice.

  Include the main piece of code in the appendix of your report, i.e. the actual realisation of the Minover update. Submit the rest of your code in separate files on Brightspace.

- ... determine the generalization error (at the end of the training process) according to

  $$\epsilon_g(t_{max}) = \frac{1}{\pi} \arccos\left(\frac{\mathbf{w} \cdot \mathbf{w}^*}{|\mathbf{w}(t_{max})| \, |\mathbf{w}^*|}\right).$$

  By repeating the training for different P, determine the so-called *learning curve*, i.e. the final generalization error $\epsilon_g(t_{max})$ as a function of the scaled number of examples $\alpha = P/N$. Obtain the result on average over $n_D \geq 10$ randomized data sets per value of P.

Consider a somewhat larger range of α than in assignment (I). The system size N and the range and number of different values of α depends, of course, on your patience, available CPU power, and the efficiency of your implementation. Provide your results for <u>at least</u> the 8 values $\alpha = 0.25, 0.5, 1.0, 1.5, 2.0, 3.0, 4.0, 5.0$. If time allows, include more intermediate values and/or extend the experiments to even larger α.

Hints:
(1) It is important that t_{max} is large enough for the algorithm to converge or at least get close to optimal stability. Run exploratory experiments first to determine a reasonable value of t_{max} to make sure that a good estimate of κ_{max} is obtained.
(2) The division by $|\mathbf{w}|$ is an important part of the definition of κ^μ. However, if you compare different κ^ν for the <u>same</u> given weight vector, i.e. when identifying the minimum, you can of course drop it. In other words: for a given $\mathbf{w}$, the minimum of the E^ν identifies the relevant example.

Suggested extensions (bonus problems):

- Define and implement a reasonable stopping criterion that ends training when the perceptron system *does not change anymore* over a (large enough) number of steps. Note that a suitable criterion is less trivial to formulate than it seems at first sight. Different criteria are possible, please provide arguments for and test your own choice.

- Determine the individual stabilities $\kappa^\nu(t_{max})$ and/or the embedding strengths x^μ at the end of training and generate histograms of the frequency of the observed values. This should be averaged over many runs of the algorithm for data sets of the same size P. All other other parameters should also be the same in all runs contributing to a histogram.

- Repeat the above experiments for the simple Rosenblatt Perceptron, which stops as soon as the classes are separated. Compare the learning curves $\epsilon_g(\alpha)$. Can you confirm that maximum stability yields better generalization behavior?

- Implement the AdaTron algorithm (in terms of embedding strengths) and compare the convergence behavior / speed with the Minover algorithm.

- Consider learning from noisy examples by replacing the true labels in the data set by

$$S^\mu = \begin{cases} +\mathrm{sign}\,(\mathbf{w}^* \cdot \boldsymbol{\xi}^\mu) & \text{with probability } 1 - \lambda \\ -\mathrm{sign}\,(\mathbf{w}^* \cdot \boldsymbol{\xi}^\mu) & \text{with probability } \lambda \,. \end{cases}$$

Here $0 < \lambda < 0.5$ controls the noise level in the training data.

- For non-separable data (e.g. with noisy labels as suggested above) implement the "large margin with error" version of the AdaTron algorithm and perform a similar set of experiments.

- **Iris Flower data**
  Determine the perceptron of maximum stability for a suitable sub-set of the Iris Flower data set. For details, see the separate document provided under "Iris data set" in the folder "Assignments".
  Methods for non-separable data could be explored for the non-separable subset of Iris data, e.g. by means of the *pocket algorithm* or the *AdaTron with errors*.